

S228/419C

DUBLIN INSTITUTE OF TECHNOLOGY
KEVIN STREET, DUBLIN 8

BSc. (Hons) in Computer Science

Stage 4

SEMESTER 2 EXAMINATIONS 2013/2014

ARTIFICIAL INTELLIGENCE II

Dr. John Kelleher
Dr. Deirdre. Lillis
Mr. P. Collins

Monday
12th May 2014
4:00 p.m to 6:00 p.m

Question 1 is **compulsory**

Answer Question 1 (40 marks) **and**
any 2 Other Questions (30 marks each).

1. (a) Explain what is meant by **inductive learning**.
(5 marks)
- (b) Inductive machine learning is often referred to as an **ill-posed problem**. What is meant by this?
(15 marks)
- (c) In the context of machine learning, explain what is meant by the term **inductive bias** and illustrate your explanation using examples of inductive biases used by machine learning algorithms.
(15 marks)
- (d) Explain what can go wrong when a machine learning classifier uses the wrong inductive bias.
(5 marks)

2. (a) A data analyst building a k -nearest neighbour model for a continuous prediction problem is considering appropriate values to use for k .
- (i) Initially the analyst uses a simple average of the target variables for the k nearest neighbours in order to make a new prediction. After experimenting with values for k in the range $0 - 10$ it occurs to the analyst that they might get very good results if they set k to the total number of instances in the training set. Do you think the analyst is likely to get good results using this value for k ?
(5 marks)
 - (ii) If the analyst was using a distance weighted averaging function rather than a simple average for their predictions would this have made their idea any more useful?
(5 marks)
- (b) A dataset showing the decisions made by an individual about whether to wait for a table at a restaurant is listed in Table 1 on the next page. (Note that Table 2, also on the next page, lists some equations that you may find useful for this question.)
- (i) Given that the *WillWait* column lists the values of the target variable, compute the entropy for this dataset.
(5 marks)
 - (ii) What is the information gain for the *Patrons* feature?
(5 marks)
 - (iii) What is the information gain for the *Type* feature?
(5 marks)
 - (iv) Given a choice between the *Patrons* and *Type* feature, which feature would the ID3 algorithm choose as the root node for a decision tree?
(5 marks)

ID	Bar	Patrons	Price	Rain	Type	WillWait
1	F	Some	\$\$\$	F	French	T
2	F	Full	\$	F	Thai	F
3	T	Some	\$	F	Burger	T
4	F	Full	\$	F	Thai	T
5	F	Full	\$\$\$	F	French	F
6	T	Some	\$\$	T	Italian	T
7	T	None	\$	T	Burger	F
8	F	Some	\$\$	T	Thai	T
9	T	Full	\$	T	Burger	F
10	T	Full	\$\$\$	F	Italian	F
11	F	None	\$	F	Thai	F
12	T	Full	\$	F	Burger	T

Table 1: A dataset describing the previous decisions made by an individual about whether to wait for a table at a restaurant.

$$\begin{aligned}
\text{Entropy}(DS) &= - \sum_{i=1}^k p_i \times \log_2(p_i) \\
\text{Remainder}(F) &= \sum_{v \in \text{Domain}(F)} \frac{|DS_v|}{|DS|} \text{Entropy}(DS_v) \\
\text{InformationGain}(F, DS) &= \text{Entropy}(DS) - \text{Remainder}(F)
\end{aligned}$$

Table 2: Equations from information theory.

3. Table 3 (on the next page) lists a dataset of the subject lines from emails. Table 4 (also on the next page) shows the subject line for an email that we would like to classify as Spam or Ham.

- (a) Using **Laplacian smoothing**, where

$$p(x = v) = \frac{\text{count}(x = v) + k}{\text{count}(x) + (k \times |\text{Domain}(x)|)}$$

with **k=1** and a **vocabulary size of 12**, calculate the following probabilities:

- (i) $P(\text{Spam}) = ?$ (2 marks)
 - (ii) $P(\text{Ham}) = ?$ (2 marks)
 - (iii) $P('Fun'|\text{Spam}) = ?$ (2 marks)
 - (iv) $P('Fun'|\text{Ham}) = ?$ (2 marks)
 - (v) $P('is'|\text{Spam}) = ?$ (2 marks)
 - (vi) $P('is'|\text{Ham}) = ?$ (2 marks)
 - (vii) $P('Free'|\text{Spam}) = ?$ (2 marks)
 - (viii) $P('Free'|\text{Ham}) = ?$ (2 marks)
- (b) Calculate the probability of the query title in Table 4 belonging to the Spam class under the **Naive Bayes assumption** and using the **smoothed probabilities** you calculated in Part (a):

$$P(\text{Spam} | 'Fun is Free') = ?$$

(7 marks)

- (c) Calculate the probability of the query title in Table 4 belonging to the Spam class under the **Naive Bayes assumption** and using **maximum likelihood** probabilities (i.e. the probabilities we could get if we did not use Laplacian smoothing):

$$P(\text{Spam} | 'Fun is Free') = ?$$

(7 marks)

Table 3: Spam and Ham Dataset

Spam	Ham
<i>Offer is Free</i>	<i>Great Learning Fun</i>
<i>Free Learning Link</i>	<i>Great Machine Learning</i>
<i>Cick Free Link</i>	<i>Free Learning Event</i>
	<i>Learning is Fun</i>
	<i>Learning Costs Money</i>

Table 4: Query Title

<i>Fun is Free</i>

x	0	1	2	3	4
y	3	6	7	8	11

Table 5: Example Dataset for Linear Regression Question

4. (a) Assuming a domain with one descriptive feature x and one target feature y , linear regression uses the following formula to model the relationship between the descriptive and target features:

$$f(x) = w_1x + w_0$$

where w_1 and w_0 are computed using the following formulae, where M is number of data points in the dataset:

$$w_1 = \frac{(M \sum_{i=1}^M x_i y_i) - (\sum_{i=1}^M x_i \sum_{i=1}^M y_i)}{(M \sum_{i=1}^M x_i^2) - (\sum_{i=1}^M x_i)^2}$$

$$w_0 = \left(\frac{1}{M} \sum_{i=1}^M y_i\right) - \left(\frac{w_1}{M} \sum_{i=1}^M x_i\right)$$

Using the data in Table 5 compute the values of w_0 and w_1 that provide the best linear fit to the data.

(10 marks)

- (b) Figure 1 (on the next pages) shows a backpropagation network that is currently processing the training vector $[1.0, 0.9, 0.9]$ which has an associated target vector $[0.1, 0.9, 0.1]$. Given that the output from unit B is 0.6 and from C is 0.8, and assuming that the activation function used at all nodes in the network is the logistic function, carry out the calculations listed below. Note that Table 6 (also on the next page) lists some equations that you may find useful when doing this question.

- (i) Calculate the actual output vector (to 3 decimal places).

(10 marks)

- (ii) Calculate the Δ error for each output unit (to 3 decimal places).

(6 marks)

- (iii) Calculate the new weight W_{BD} for the connection from unit B to the output unit D after the training example has been processed. Use a learning rate of $\eta = 0.25$.

(4 marks)

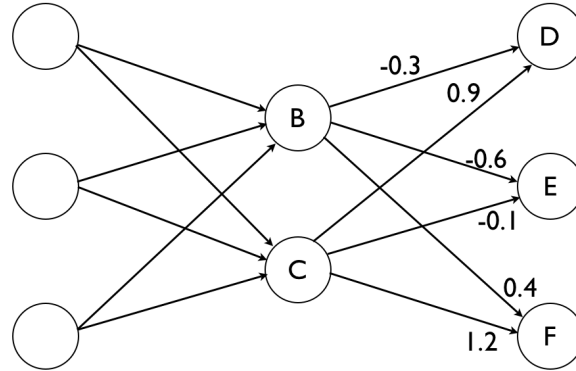


Figure 1: Example Neural Net

Weighted sum of inputs for unit i with j inputs:	$in_i = \sum_j W_{ji} a_j(in_j)$
Activation Function (Logistic) for unit i :	$a_i(in_i) = \frac{1}{1 + \exp^{-in_i}}$
Perceptron weight update rule for link $j \rightarrow i$	$w_{ji} = w_{ji} + \eta (t_i - a_i(in_i)) \times a_j(in_j)$
Hebbian Weight Update Rule for link $j \rightarrow i$	$w_{ji} = \eta \times a_j(in_j) \times a_i(in_i)$
Partial Derivative for Logistic Activation Function	$\frac{\delta a_i(in_i)}{\delta in_i} = a_i(in_i) \times (1 - a_i(in_i))$
Error for an output unit i	$error_i = target_i - a_i(in_i)$
Delta Error for an output unit i	$\Delta_i = error_i \times a_i(in_i) \times (1 - a_i(in_i))$
Delta Error for a hidden unit j feeding into n units	$\Delta_j = \left(\sum_{i=1}^n W_{ji} \times \Delta_i \right) \times a_j(in_j) \times (1 - a_j(in_j))$
Delta Weight Update Rule for link $x \rightarrow k$	$W_{x,k} = W_{x,k} + (\eta \times a_x(in_x) \times \Delta_k)$

Table 6: Equations used in Perceptron and Neural Network training.